

TEXTBOOK EDITION • PROBLEM SOLVING & ALGORITHMS

Data Structures & Algorithms (FinTech Patterns)

Arrays, Two Pointers, Sliding Window, Hashing & Algorithmic Foundations

TARGET ROLE	Developer (3+ Years Experience)
DEPARTMENT	Strategic Projects & New Age Tech, Bengaluru
PREPARED FOR	Bipin Yadav (SDE II, 4+ YoE)
ESTIMATED STUDY TIME	18 Hours
DIFFICULTY LEVEL	Intermediate to Advanced
TEACHING FORMAT	Hinglish Mental Models + Production FinTech

Table of Contents

Chapter 01 — Algorithmic Complexity Analysis (Big-O, Big-Ω, Big-Θ & Amortized Analysis)

Study Time: 35m

Chapter 02 — Array Patterns, In-Place Manipulation & Prefix Sums

Study Time: 35m

Chapter 03 — Hashing Patterns, Hash Tables & Frequency Counting

Study Time: 35m

Chapter 04 — Two Pointers Algorithmic Patterns

Study Time: 35m

Chapter 05 — Sliding Window Algorithmic Patterns

Study Time: 35m

Chapter 06 — Python for Coding Interviews & DSA Mastery

Study Time: 35m

Chapter 07 — DSA Last-Minute Algorithmic Patterns Cheatsheet

Study Time: 35m

CHAPTER 01

Algorithmic Complexity Analysis (Big-O, Big-Ω, Big-Θ & Amortized Analysis)

PRIORITY

 High Priority

DIFFICULTY

Intermediate to Advanced

EST. TIME

35–45 min

FREQUENCY

Very High (85%+)

Algorithmic Complexity Analysis (Big-O, Big-Ω, Big-Θ & Amortized Analysis)

1. Why This Matters

In banking and financial backend systems, an algorithm that is $O(N^2)$ instead of $O(N \log N)$ can freeze a settlement engine during peak trading or UPI clearing windows. At IDFC FIRST Bank's Strategic Projects, engineers write code handling tens of thousands of concurrent payment requests. In interviews, stating an algorithm's complexity is not enough: interviewers will ask you to mathematically derive worst-case, average-case, and amortized bounds, identify auxiliary memory allocations, and analyze garbage collection overhead.

2. Prerequisites

- Basic algebra (logarithms, exponents, summation formulas $\sum_{i=1}^n i = \frac{n(n+1)}{2}$).
- Understanding of stack frames and heap allocations in programming languages (Node.js/V8, Python).

3. Concept

Asymptotic Notations

- Big-O (O):** Upper bound. Formal definition: $f(N) = O(g(N))$ if there exist positive constants c and N_0 such that $0 \leq f(N) \leq c \cdot g(N)$ for all $N \geq N_0$. Represents the worst-case scenario.
- Big-Omega (Ω):** Lower bound. $f(N) \geq c \cdot g(N)$ for all $N \geq N_0$. Best-case performance.
- Big-Theta (Θ):** Tight bound. When an algorithm's upper bound and lower bound match: $c_1 \cdot g(N) \leq f(N) \leq c_2 \cdot g(N)$.
- Little-o (o):** Strict upper bound where $f(N)$ grows strictly slower than $g(N)$.

```

flowchart LR
    subgraph Growth["Growth of Common Time Complexities"]
        direction TB
        c1["O(1) - Constant: Hash table lookup, Array indexing"]
        c2["O(log N) - Logarithmic: Binary search, Balanced BST lookup"]
        c3["O(N) - Linear: Single scan, Two Pointers, Sliding Window"]
        c4["O(N log N) - Linearithmic: MergeSort, QuickSort (avg), Timsort"]
        c5["O(N^2) - Quadratic: Nested loops, BubbleSort, Brute force pairs"]
        c6["O(2^N) - Exponential: Naive Fibonacci, Subsets recursion"]
        c7["O(N!) - Factorial: Permutations generation, Traveling Salesperson"]
    end
    c1 --> c2 --> c3 --> c4 --> c5 --> c6 --> c7

```

Amortized Analysis

When an operation is occasionally very expensive but predominantly cheap, the *amortized cost* averages the total cost over a long sequence of operations. **The Accounting Method:** You overcharge cheap operations (e.g., normal array append = 1 unit of work charged as 3 units) and store the surplus as credit. When the expensive resizing operation occurs (copying N elements = N units of work), the accumulated credit pays for the reallocation. Thus, the amortized cost per append is strictly $O(1)$.

4. Simple Example: Linear vs Logarithmic Search

```

# Linear Search: O(N) Time, O(1) Space
def linear_search(arr: list[int], target: int) → int:
    for i in range(len(arr)):
        if arr[i] == target:
            return i
    return -1

# Binary Search: O(log N) Time, O(1) Space
def binary_search(sorted_arr: list[int], target: int) → int:
    left, right = 0, len(sorted_arr) - 1
    while left ≤ right:
        mid = left + (right - left) // 2 # Prevents integer overflow
        if sorted_arr[mid] == target:
            return mid
        elif sorted_arr[mid] < target:
            left = mid + 1
        else:
            right = mid - 1
    return -1

```

5. Real-World Banking Example: High-Throughput Transaction Deduplication

Suppose an incoming UPI payment batch has $N = 100,000$ transactions. You must verify that no two transactions share the same `idempotency_key`.

- **Brute Force ($O(N^2)$):** Comparing each transaction against every other transaction requires:
 $\frac{100,000 \times 99,999}{2} \approx 5 \times 10^9$ operations $\implies \approx 5 \times 10$ seconds (CPU freeze).

- **Hash Set Optimization ($O(N)$ Time, $O(N)$ Space):** Checking existence in a hash set takes $O(1)$ on average: $\$100,000 \text{ lookups} \approx 0.015 \text{ seconds}$.
- **In-place Sort + Adjacent Scan ($O(N \log N)$ Time, $O(1)$ Auxiliary Space):** Useful when memory is constrained on embedded gateway appliances.

6. Code: Python Dynamic Array Resizing Proof

```
import sys
import time

def demonstrate_amortized_growth():
    """Demonstrates list capacity jumps in Python (CPython dynamic array)."""
    data = []
    prev_size = sys.getsizeof(data)
    print(f"{'Length':<10} | {'Byte Size':<12} | {'New Allocation Triggered?'}")
    print("-" * 50)

    for i in range(25):
        data.append(i)
        current_size = sys.getsizeof(data)
        resized = current_size != prev_size
        if resized:
            print(f"{'len(data)':<10} | {'current_size':<12} | YES (Grown to accommodate elements)")
            prev_size = current_size

if __name__ == "__main__":
    demonstrate_amortized_growth()
```

7. How It Works Internally

CPython & V8 Dynamic Array Growth Mechanics

1. In CPython, `list` is implemented as an array of pointers (`PyObject**`). When the allocated capacity is exhausted, CPython over-allocates using the formula: $\text{new_allocated} = \text{size} + (\text{size} // 3) + (\text{size} < 9 ? 3 : 6)$. This provides an approximate geometric growth factor of $\approx 1.125 \times$.
2. In V8 (Node.js/JavaScript), `JSArray` has two modes: **Fast Elements** (contiguous C++ pointer backing stores) and **Dictionary Elements** (hash tables for sparse arrays). For fast elements, resizing multiplies capacity by $1.5 \times$ to $2 \times$.
3. Copying N elements takes $O(N)$, but because reallocation happens geometrically, resizing from 1 to N requires: $1 + 2 + 4 + 8 + \dots + N = 2N - 1 = O(N)$ total copies across N insertions. $\frac{\text{Total Time}}{\text{Operations}} = \frac{O(N)}{N} = O(1)$ amortized per append.

8. Common Mistakes

- Confusing Auxiliary Space with Total Space:** Auxiliary space is temporary working memory *excluding* the input. In MergeSort, total space is $O(N)$, and auxiliary space is $O(N)$. In QuickSort, auxiliary stack space is $O(\log N)$ average, $O(N)$ worst.
- Ignoring String Immutability Overhead:** In Python and JavaScript, strings are immutable. Doing `s += char` inside an N -iteration loop creates a new string each time, turning an intended $O(N)$ operation into $O(N^2)$! Always use `"".join(list)` in Python or an array push + `.join("")` in JS.
- Assuming Hash Table Lookups are Always $O(1)$:** If hash collisions occur (or an adversary triggers hash collision attacks), open-addressed or chained tables degrade to $O(N)$ worst case.

9. Performance / Complexity Matrix

Operation / Algorithm	Best Time	Average Time	Worst Time	Auxiliary Space
Array Access by Index	$O(1)$	$O(1)$	$O(1)$	$O(1)$
Dynamic Array Append	$O(1)$	$O(1)$	$O(N)$	$O(1)$
Hash Table Insert/Lookup	$O(1)$	$O(1)$	$O(N)$	$O(N)$
Binary Search	$O(1)$	$O(\log N)$	$O(\log N)$	$O(1)$
Merge Sort	$O(N \log N)$	$O(N \log N)$	$O(N \log N)$	$O(N)$
Quick Sort	$O(N \log N)$	$O(N \log N)$	$O(N^2)$	$O(\log N)$ avg
Breadth-First Search (BFS)	$O(V + E)$	$O(V + E)$	$O(V + E)$	$O(V)$
Depth-First Search (DFS)	$O(V + E)$	$O(V + E)$	$O(V + E)$	$O(V)$ stack

10. Interview Questions (Easy to Medium to Hard)

Easy

- Q:** What is the difference between $O(1)$ and $O(\log N)$? Give an example of each in software engineering.

Medium

- Q:** Explain why QuickSort is commonly preferred over MergeSort for arrays in cache-sensitive systems despite QuickSort having a worse worst-case time complexity ($O(N^2)$ vs $O(N \log N)$).

Hard

- **Q:** An algorithm processes a stream of financial transactions. Each incoming item requires checking a min-heap of size K and occasionally rebuilding an auxiliary index of size M . Derive the tight asymptotic bound for N total transactions where $M = \sqrt{N}$ and $K = \log N$.

11. Follow-up Questions from Interviewer

- "What is cache locality, and how does it affect the empirical execution speed of an $O(N)$ array traversal versus an $O(N)$ linked list traversal?"
- "If `Array.prototype.sort()` in Node.js uses Timsort, what is its worst-case space complexity, and why doesn't it use standard in-place QuickSort?"

12. Model Answer (QuickSort vs MergeSort & Cache Locality)

BANKING & FINTECH ARCHITECTURE SCENARIO

Interviewer: "Why does production software often choose QuickSort or Dual-Pivot QuickSort over MergeSort for primitive arrays?"

Model Answer: "While MergeSort guarantees $O(N \log N)$ worst-case time, it requires $O(N)$ auxiliary memory because merging requires a secondary array. In contrast, QuickSort operates strictly in-place with $O(\log N)$ auxiliary call-stack space.

More importantly, QuickSort exhibits superior **spatial locality of reference**. Its partitioning algorithm traverses contiguous memory sequentially, making optimal use of CPU L1/L2 cache lines and minimizing hardware cache misses. MergeSort's auxiliary buffer copying incurs additional memory bus traffic.

However, for reference types or objects (such as JavaScript arrays or Python objects) where elements are pointers and stability is required (preserving relative order of equal keys, like sorting bank transactions by timestamp while preserving category order), Timsort (a hybrid of MergeSort and InsertionSort) is preferred because it guarantees stability and $O(N)$ best-case time on partially sorted data."

13. Practical Exercise

Analyze the time and auxiliary space complexity of each function:

```
# Exercise 1:
def mystery_one(n: int):
    count = 0
    i = n
    while i > 0:
        for j in range(i):
            count += 1
        i //= 2
    return count

# Exercise 2:
def mystery_two(arr: list[int]):
    n = len(arr)
    result = []
    for i in range(n):
        subset = []
        for j in range(i, n):
            subset.append(arr[j])
            result.append(list(subset))
    return result
```

(Solution hint: Mystery 1 uses a geometric series $\sum_{k=0}^{\log_2 n} \frac{n}{2^k} = n(1 + 1/2 + 1/4 + \dots) = O(N)$ time, $O(1)$ space. Mystery 2 generates all subarrays: $O(N^2)$ iterations, but copying subsets makes total time $O(N^3)$ and space $O(N^3)$).

14. Quick Revision

- Big-O denotes worst-case upper bound; Big- Θ denotes tight asymptotic bound.
- Geometric series $N + N/2 + N/4 + \dots = 2N = O(N)$.
- In-place algorithms utilize $O(1)$ auxiliary space beyond input buffers.
- String concatenation inside loops in Python/JS creates $O(N^2)$ memory churning.
- Dynamic array append is $O(1)$ amortized, $O(N)$ worst-case.

15. Interview Checklist

- ☐ Can define Big-O mathematically ($f(N) \leq c \cdot g(N)$).
- ☐ Can explain the accounting/potential method for amortized analysis.
- ☐ Can distinguish auxiliary space from input space.
- ☐ Can explain CPU cache lines and spatial/temporal locality.
- ☐ Can identify hidden $O(N)$ costs in high-level language methods (e.g., `list.insert(0, val)` or `Array.shift()` which shifts all elements).

 CHAPTER 01 SUMMARY & REVISION CHECKLIST **WHAT TO REMEMBER**

- Basic algebra (logarithms, exponents, summation formulas $\sum_{i=1}^n i = \frac{n(n+1)}{2}$).
- Understanding of stack frames and heap allocations in programming languages (Node.js/V8, Python).
- Big-O (O): Upper bound. Formal definition: $f(N) = O(g(N))$ if there exist positive constants c and N_0 such that

 CANDIDATE SELF-CHECK

- ☐ I understand the fundamental mechanism
- ☐ I can explain it in natural Hinglish
- ☐ I can write or debug the core code/query
- ☐ I can defend trade-offs and edge cases

CHAPTER 02

Array Patterns, In-Place Manipulation & Prefix Sums

PRIORITY

● High Priority

DIFFICULTY

Intermediate to Advanced

EST. TIME

35–45 min

FREQUENCY

Very High (85%+)

Array Patterns, In-Place Manipulation & Prefix Sums

1. Why This Matters

Arrays form the backbone of memory management and algorithmic screening at IDFC FIRST Bank. In live coding interviews, array problems test your ability to avoid auxiliary allocations, identify prefix accumulations, manipulate boundaries without off-by-one errors, and optimize brute-force $O(N^2)$ loops into $O(N)$ linear scans.

2. Prerequisites

- Familiarity with zero-indexed array indexing and contiguous memory buffers.
- Understanding of time vs space complexity trade-offs.

3. Concept

Key Array Algorithmic Patterns

1. **Prefix Sum (Cumulative Sum):** Precomputing running totals allows any range sum query `sum(arr[i...j])` to be answered in $O(1)$ time via `prefix[j+1] - prefix[i]`.
2. **In-Place Two Pointers / Swapping:** Modifying array values without creating auxiliary lists, maintaining $O(1)$ extra space.
3. **Dutch National Flag (Three-Way Partitioning):** Partitioning an array around pivot values in a single linear pass.

```
flowchart LR
    subgraph PrefixSum["Prefix Sum Array Construction"]
        arr["arr: [ 10, 20, 15, 30 ]"]
        pre["pref: [ 0, 10, 30, 45, 75 ]"]
    end
    end
    arr --> pre
```

$$\text{RangeSum}(i, j) = \text{pref}[j+1] - \text{pref}[i]$$

4. Simple Example: In-Place Removal of Duplicates

```
def remove_duplicates_sorted(nums: list[int]) → int:
    """
    Removes duplicates in-place from a sorted array.
    Returns the count of unique elements.
    Complexity: O(N) Time, O(1) Auxiliary Space.
    """
    if not nums:
        return 0

    write_ptr = 1
    for read_ptr in range(1, len(nums)):
        if nums[read_ptr] != nums[read_ptr - 1]:
            nums[write_ptr] = nums[read_ptr]
            write_ptr += 1

    return write_ptr
```

5. Real-World Banking Example: Daily Transaction Range Sum Engine

A core banking ledger receives thousands of audit requests: "What was the total volume transacted between timestamp T_1 and T_2 ?" - Re-summing transactions on each query is $O(N)$ per request, overloading the database under high concurrency. - By maintaining an hourly or daily **Prefix Sum Ledger**, the balance query computes in $O(1)$ instant time via a simple subtraction.

6. Code: Product of Array Except Self ($O(N)$ Time, $O(1)$ Extra Space)

Problem Statement

Given an integer array `nums`, return an array `answer` such that `answer[i]` is equal to the product of all the elements of `nums` except `nums[i]`. You must solve it **without using the division operation** and in $O(N)$ time.

Clarifying Questions

1. Can elements be zero or negative? Yes, zeroes are explicitly allowed.
2. Can the product overflow a standard 32-bit integer? Yes, in Python integers have arbitrary precision, but in languages like Java/C++, 64-bit integers (`long`) must be used.
3. Does the output array count toward auxiliary space? No, output space is excluded from auxiliary space analysis.

Brute-Force Approach

For every index i , iterate through the rest of the array with a nested loop and calculate the product. -

Time Complexity: $O(N^2)$ - **Space Complexity:** $O(1)$

Optimized Approach: Prefix and Suffix Running Products

Instead of recomputing products repeatedly, notice that: $\text{product_except_self}[i] = (\text{product of elements to the left of } i) \times (\text{product of elements to the right of } i)$

```
def product_except_self(nums: list[int]) → list[int]:
    n = len(nums)
    output = [1] * n

    # Step 1: Compute prefix products in-place in output array
    prefix = 1
    for i in range(n):
        output[i] = prefix
        prefix *= nums[i]

    # Step 2: Traverse backwards and multiply by running suffix product
    suffix = 1
    for i in range(n - 1, -1, -1):
        output[i] *= suffix
        suffix *= nums[i]

    return output

# Test
print(product_except_self([1, 2, 3, 4])) # [24, 12, 8, 6]
print(product_except_self([-1, 1, 0, -3, 3])) # [0, 0, 9, 0, 0]
```

- **Time Complexity:** $O(N)$ (Two passes of length N).
- **Space Complexity:** $O(1)$ auxiliary space (only scalar `prefix` and `suffix` trackers).
- **Edge Cases:** Array with a single zero, multiple zeroes (where all outputs become zero), all negative numbers.

7. How It Works Internally: Memory Layout & Cache Locality

Arrays in RAM are stored in a contiguous block of virtual memory. When the CPU executes `output[i] = prefix`, hardware prefetchers load a **cache line** (typically 64 bytes) into L1/L2 cache. – Forward and backward sequential scans have **100% spatial locality**, meaning almost zero cache misses. – If this were implemented using a linked list or pointer-heavy tree, each node access would require an unpredictable memory hop, causing severe cache line stalls.

8. Common Mistakes

1. **Division by Zero:** Attempting to solve *Product of Array Except Self* by calculating the total product and dividing by `nums[i]` crashes when `nums[i] == 0`.
2. **Off-by-One in Prefix Sums:** Creating a prefix array of size N instead of $N+1$. Using size $N+1$ with `prefix[0] = 0` eliminates messy `if i == 0:` boundary checks.
3. **Modifying Array While Iterating:** Mutating an array with `.pop()` or `.insert()` during a `for x in arr:` loop alters the iteration index, leading to skipped elements.

9. Performance / Complexity Matrix

Problem / Technique	Brute Force	Optimized Time	Auxiliary Space
Prefix Sum Range Query	$O(N)$ per query	$O(1)$ per query ($O(N)$ precompute)	$O(N)$
Product Except Self	$O(N^2)$	$O(N)$	$O(1)$
Remove Duplicates In-Place	$O(N^2)$	$O(N)$	$O(1)$
Rotate Array by K steps	$O(N \times K)$	$O(N)$ (Triple reverse)	$O(1)$

10. Interview Questions (Easy $\rightarrow$ Medium $\rightarrow$ Hard)

Easy

- **Q:** How do you reverse an array in-place without using extra memory?

Medium

- **Q:** Given an unsorted array of integers containing positive and negative numbers, find the contiguous subarray which has the largest sum (Kadane's Algorithm).

Hard

- **Q:** Given an array representing stock prices across consecutive trading ticks, find the maximum profit you can achieve with at most two transactions.

11. Follow-up Questions from Interviewer

- "In Kadane's algorithm, what happens if all numbers in the array are negative? How do you ensure your code doesn't mistakenly return 0?"
- "Can you modify your range sum query to handle dynamic in-place updates in $O(\log N)$ time? (Segment Tree / Fenwick Tree)"

12. Model Answer: Kadane's Algorithm ($O(N)$ Maximum Subarray)

30-SECOND INTERVIEW ANSWER (EXECUTIVE PITCH)

Interviewer: "Explain how Kadane's algorithm works and prove its optimality."

Model Answer: "Kadane's algorithm uses dynamic programming with state reduction to find the maximum sum contiguous subarray in $O(N)$ time and $O(1)$ auxiliary space."

At each index i , we make a local decision: Is it better to extend the existing contiguous subarray ending at $i-1$, or start a brand new subarray beginning at i ?

The recurrence relation is: $\text{current_max} = \max(\text{nums}[i], \text{current_max} + \text{nums}[i])$

If `current_max` ever becomes negative, adding it to the next element would strictly decrease that element's potential sum. Therefore, we discard the negative running sum and restart at the current element.

To properly handle arrays containing all negative numbers (e.g., `[-5, -2, -8]`), we initialize `max_so_far` and `current_max` to `nums[0]` rather than `0`. This guarantees we return the least negative number (`-2`) rather than an invalid empty subarray."

```
def max_sub_array(nums: list[int]) → int:
    current_max = nums[0]
    global_max = nums[0]
    for num in nums[1:]:
        current_max = max(num, current_max + num)
        global_max = max(global_max, current_max)
    return global_max
```

13. Practical Exercise

Implement the **Rotate Array by K Steps** in-place in $O(N)$ time and $O(1)$ space using the 3-step array reversal technique: 1. Reverse entire array. 2. Reverse first K elements. 3. Reverse remaining $N - K$ elements.

(Ensure $K = K \bmod N$ handles cases where $K \geq N$).

14. Quick Revision

- Prefix sum: $\text{sum}(i, j) = \text{pref}[j+1] - \text{pref}[i]$.
- In-place two pointers: read pointer scans ahead, write pointer retains valid output.
- Kadane's: $\text{local} = \max(\text{num}, \text{local} + \text{num})$.
- In-place reversal eliminates auxiliary space for rotation problems.

15. Interview Checklist

- [] Writes Kadane's algorithm correctly for all-negative inputs.

- [] Solves Product Except Self without division.
- [] Understands why $N+1$ sized prefix arrays prevent boundary condition bugs.
- [] Explains cache line prefetching benefits of contiguous memory.

CHAPTER 02 SUMMARY & REVISION CHECKLIST

WHAT TO REMEMBER

- Familiarity with zero-indexed array indexing and contiguous memory buffers.
- Understanding of time vs space complexity trade-offs.
- Re-summing transactions on each query is $O(N)$ per request, overloading the database under high concurrency.

CANDIDATE SELF-CHECK

- ☐ I understand the fundamental mechanism
- ☐ I can explain it in natural Hinglish
- ☐ I can write or debug the core code/query
- ☐ I can defend trade-offs and edge cases

CHAPTER 03

Hashing Patterns, Hash Tables & Frequency Counting

PRIORITY

● High Priority

DIFFICULTY

Intermediate to Advanced

EST. TIME

35–45 min

FREQUENCY

Very High (85%+)

Hashing Patterns, Hash Tables & Frequency Counting

1. Why This Matters

Hash tables are the most ubiquitous data structure in modern backend engineering. At IDFC FIRST Bank, hash maps power distributed session caching, idempotency tracking, fraud velocity checks, and fast in-memory record matching. In coding interviews, recognizing that an $O(N^2)$ brute-force lookup can be converted into an $O(N)$ linear time algorithm using a hash map or hash set is the single most frequent optimization pattern.

2. Prerequisites

- Understanding of key-value associations.
- Basic knowledge of hash functions and bucket indexing.

3. Concept

Hash Map Fundamentals

A hash table maps keys to values by passing the key through a **hash function** to compute an integer hash code, which is then mapped to an internal bucket index using modulo arithmetic: $\text{bucket_index} = \text{hash}(\text{key}) \bmod \text{capacity}$

```
flowchart LR
    subgraph Hashing["Hash Function & Bucket Resolution"]
        k1["Key: 'pan_ABCDE1234F'"] --> h["Hash Function"]
        h --> code["Hash: 0x8F3A29"]
        code --> mod["Modulo Capacity (16)"]
        mod --> b9["Bucket #9"]
    end
```

Collision Resolution Strategies

1. **Separate Chaining:** Each bucket stores a linked list (or balanced red-black tree, as in Java 8+ `HashMap`) of key-value pairs that hash to the same bucket.

2. **Open Addressing (Linear/Quadratic/Pseudo-random Probing)**: All entries reside directly within the bucket array. If a collision occurs, the algorithm systematically probes adjacent or pseudo-random indices until an empty slot is found (used by Python dictionaries).

4. Simple Example: Two Sum ($O(N)$ Time, $O(N)$ Space)

```
def two_sum(nums: list[int], target: int) → list[int]:
    """
    Finds indices of two numbers that add up to target.
    Complexity:  $O(N)$  Time,  $O(N)$  Space.
    """
    seen = {} # value → index
    for i, num in enumerate(nums):
        complement = target - num
        if complement in seen:
            return [seen[complement], i]
        seen[num] = i
    return []
```

5. Real-World Banking Example: Real-Time Fraud Velocity Detection

During a flash sale, card fraud detection rules require checking: "Has this card transacted more than 3 times in the past 60 seconds?" - Scanning the relational transactions table each time would overwhelm the database with disk I/O. - By maintaining an in-memory hash table (or Redis sorted set / hash), the system can record timestamped attempts by `card_id` and check velocity thresholds in sub-millisecond $O(1)$ time.

6. Code: Subarray Sum Equals K (Prefix Sum + Hash Map)

Problem Statement

Given an array of integers `nums` and an integer `k`, return the total number of continuous subarrays whose sum equals to `k`.

Clarifying Questions

1. *Can elements be negative?* Yes! Because numbers can be negative, a sliding window **cannot** be used because the running sum does not increase monotonically.
2. *Can the array be empty?* Assume $1 \leq \text{len} \leq 20,000$.

Intuition: Prefix Sum Complement

Let $\text{pref}[i]$ be the sum of elements from index 0 to i . A subarray from j to i has sum equal to k if and only if: $\text{pref}[i] - \text{pref}[j-1] = k \iff \text{pref}[j-1] = \text{pref}[i] - k$

By storing the frequency of all previous prefix sums in a hash map, we can instantly look up how many times $\text{pref}[i] - k$ has appeared.

```

from collections import defaultdict

def subarray_sum_k(nums: list[int], k: int) → int:
    prefix_counts = defaultdict(int)
    prefix_counts[0] = 1 # Base case: a prefix sum of 0 has occurred once

    running_sum = 0
    total_subarrays = 0

    for num in nums:
        running_sum += num
        # Check how many prior prefixes satisfy: running_sum - prior_prefix = k
        complement = running_sum - k
        if complement in prefix_counts:
            total_subarrays += prefix_counts[complement]

        prefix_counts[running_sum] += 1

    return total_subarrays

# Test cases
print(subarray_sum_k([1, 1, 1], 2)) # 2 ([1,1] at index 0..1 and 1..2)
print(subarray_sum_k([1, -1, 0], 0)) # 3 ([1,-1], [0], [1,-1,0])
print(subarray_sum_k([3, 4, 7, 2, -3, 1, 4, 2], 7)) # 4

```

- **Time Complexity:** $O(N)$ single pass.
- **Space Complexity:** $O(N)$ to store prefix frequencies in the hash map.
- **Edge Cases:** Array with zeroes, negative numbers, $k = 0$, single element matching k .

7. How It Works Internally: The Hash Collision Vulnerability (HashDoS)

If an attacker knows the hashing algorithm used by an application server (e.g., standard SipHash or MurmurHash without a random seed), they can send thousands of HTTP POST parameters designed to produce identical bucket indices. - When thousands of keys collide in the same bucket, lookup complexity degrades from $O(1)$ to $O(N)$. - The server CPU spikes to 100% parsing a single malicious request. Modern runtimes (Node.js/V8, Python 3.4+) randomize the hash seed (`PYTHONHASHSEED`) on startup to prevent HashDoS attacks.

8. Common Mistakes

1. **Forgetting the Base Case** `prefix_counts[0] = 1` : In prefix sum hashing problems, omitting `prefix_counts[0] = 1` causes the algorithm to miss any valid subarray that begins at index 0.
2. **Using Mutable Keys in Hash Maps:** Keys in a hash map must be immutable (hashable). In Python, lists cannot be keys (`TypeError: unhashable type: 'list'`); use a `tuple` instead.
3. **Assuming Hash Maps are Ordered Across All Languages:** While Python 3.7+ and JavaScript maintain insertion order, standard hash maps in C++ (`std::unordered_map`) or Go maps do not guarantee order. Never rely on order unless language-guaranteed.

9. Performance / Complexity Matrix

Problem / Technique	Brute Force Time	Hash Map Optimized Time	Auxiliary Space
Two Sum	$O(N^2)$	$O(N)$	$O(N)$
Subarray Sum Equals K	$O(N^2)$	$O(N)$	$O(N)$
Group Anagrams	$O(N \cdot K \cdot \log K)$	$O(N \cdot K)$ (Character count key)	$O(N \cdot K)$
Longest Consecutive Sequence	$O(N \log N)$ (Sort)	$O(N)$ (Hash Set)	$O(N)$

10. Interview Questions (Easy to Medium to Hard)

Easy

- Q: How does a Hash Set differ from a Hash Map under the hood?

Medium

- Q: Given an unsorted array of integers, find the length of the longest consecutive elements sequence in $O(N)$ time.

Hard

- Q: Design an in-memory Key-Value store with `set(key, val)`, `get(key)`, `delete(key)`, and `getRandomKey()` all operating in strictly $O(1)$ worst-case or average time.

11. Follow-up Questions from Interviewer

- "In the $O(1)$ Key-Value store with `getRandomKey()`, why is a hash map alone insufficient, and why must you combine a dynamic array with a hash map?"
- "How does Java 8's HashMap handle high collision buckets differently than Python's open-addressing dictionary?"

12. Model Answer: Longest Consecutive Sequence in $O(N)$ Time

30-SECOND INTERVIEW ANSWER (EXECUTIVE PITCH)

Interviewer: "Given `[100, 4, 200, 1, 3, 2]`, how do you find the longest consecutive sequence (`[1, 2, 3, 4]` of length 4) in $O(N)$ without sorting?"

Model Answer: "Sorting takes $O(N \log N)$, so we must use a Hash Set for $O(1)$ lookups.

The core insight is identifying the **start of a sequence**. A number x can only be the start of a streak if $x - 1$ is NOT present in the set.

1. Insert all numbers into a Python `set`.
2. Iterate through each number in the set. If $x - 1 \in \text{set}$, we skip it because it is part of an existing streak that will be counted from its true beginning.
3. If $x - 1 \notin \text{set}$, we have found the sequence origin. We then enter an inner `while` loop checking for $x+1, x+2, \dots$ until the streak breaks.

Although there is a nested `while` loop, each number is visited at most twice: once during the set iteration, and at most once during a streak expansion. Therefore, the total time complexity is strictly $O(N)$, with $O(N)$ auxiliary space for the set."

```
def longest_consecutive(nums: list[int]) → int:
    num_set = set(nums)
    longest_streak = 0

    for num in num_set:
        # Check if num is the start of a streak
        if num - 1 not in num_set:
            current_num = num
            current_streak = 1

            while current_num + 1 in num_set:
                current_num += 1
                current_streak += 1

            longest_streak = max(longest_streak, current_streak)

    return longest_streak
```

13. Practical Exercise

Implement **Group Anagrams** without sorting strings (using a 26-element tuple of character frequencies as the hash key) to achieve $O(N \cdot K)$ time complexity where N is word count and K is max word length.

14. Quick Revision

- Hash map lookup/insert is $O(1)$ average, $O(N)$ worst-case during severe collisions.
- Subarray sum problems with negative numbers require Prefix Sum + Hash Map.

- Keys must be immutable; mutable keys corrupt bucket invariants.
- Python uses open addressing; Java uses separate chaining (converting linked list to red-black tree when bucket size exceeds 8).

15. Interview Checklist

- ☐ Remembers `prefix_counts[0] = 1` for prefix sum hashing problems.
- ☐ Understands why sliding window fails when negative numbers are present.
- ☐ Can explain how to combine an array and a hash map for $O(1)$ random access deletion.
- ☐ Explains hash seed randomization to mitigate HashDoS.



CHAPTER 03 SUMMARY & REVISION CHECKLIST



WHAT TO REMEMBER

- Understanding of key-value associations.
- Basic knowledge of hash functions and bucket indexing.
- Scanning the relational transactions table each time would overwhelm the database with disk I/O.



CANDIDATE SELF-CHECK

- ☐ I understand the fundamental mechanism
- ☐ I can explain it in natural Hinglish
- ☐ I can write or debug the core code/query
- ☐ I can defend trade-offs and edge cases

CHAPTER 04

Two Pointers Algorithmic Patterns

PRIORITY

● High Priority

DIFFICULTY

Intermediate to Advanced

EST. TIME

35–45 min

FREQUENCY

Very High (85%+)

Two Pointers Algorithmic Patterns

1. Why This Matters

Two-pointer techniques are fundamental to passing algorithmic screening interviews at top engineering hubs like IDFC FIRST Bank. The pattern eliminates nested loops ($O(N^2)$) and reduces problems to linear $O(N)$ scans with $O(1)$ auxiliary space. Interviewers frequently test two pointers on financial series, sorted settlement logs, deduplication, and pair-matching scenarios.

2. Prerequisites

- Zero-indexed array addressing.
- Sorting fundamentals (understanding that pre-sorting an array takes $O(N \log N)$ and often unlocks $O(N)$ two-pointer traversal).

3. Concept

Taxonomy of Two Pointers

1. **Opposite-Direction Pointers (Converging):** One pointer starts at the beginning (`left = 0`) and the other at the end (`right = len - 1`). They move toward each other based on comparison conditions (e.g., Two Sum on sorted array, 3-Sum, Container With Most Water).
2. **Same-Direction Pointers (Fast & Slow):** Both pointers start at index 0 , but advance at different rates or under different conditions (e.g., Linked List cycle detection, in-place duplicate removal).
3. **Dual-Array Pointers:** One pointer in Array A and one pointer in Array B (e.g., Merge Sorted Arrays, intersection of two banking transaction logs).

```
flowchart LR
  subgraph Converging["Opposite-Direction (Sorted Array)"]
    L["Left Pointer →"] --> E["[ 2, 5, 7, 11, 15 ]"]
    R["← Right Pointer"] --> E
  end
```

4. Simple Example: Two Sum on a Sorted Array ($O(N)$ Time, $O(1)$ Space)

```
def two_sum_sorted(nums: list[int], target: int) → list[int]:
    """
    Finds 1-based indices of two numbers that sum to target.
    Requires array to be pre-sorted.
    """
    left, right = 0, len(nums) - 1
    while left < right:
        current_sum = nums[left] + nums[right]
        if current_sum == target:
            return [left + 1, right + 1]
        elif current_sum < target:
            left += 1 # Need larger sum
        else:
            right -= 1 # Need smaller sum
    return []
```

5. Real-World Banking Example: Inter-Bank Netting & Settlement Matching

At end-of-day clearing, the banking switch matches credit claims against debit obligations across institutions. Given sorted lists of receivables and payables, dual pointers can reconcile mutual netting obligations in a single $O(N + M)$ pass without building multi-gigabyte hash tables in memory.

6. Code: 3-Sum (Finding Unique Triplets Summing to Zero)

Problem Statement

Given an integer array `nums`, return all the triplets `[nums[i], nums[j], nums[k]]` such that $i \neq j, i \neq k, j \neq k$, and $\text{nums}[i] + \text{nums}[j] + \text{nums}[k] == 0$. The solution set must **not contain duplicate triplets**.

Clarifying Questions

1. Can the elements be negative? Yes.
2. Does the order of triplets in the output matter? No.
3. What should be returned if fewer than 3 elements exist? Return an empty list `[]`.

Optimized Algorithm: Sort + Converging Two Pointers

1. Sort the array in $O(N \log N)$.
2. Iterate `i` from 0 to $N - 3$. If $\text{nums}[i] > 0$, break early (three positive numbers can never sum to 0).
3. Skip duplicate values of `nums[i]` to prevent duplicate triplets.
4. For each fixed `i`, initialize `left = i + 1` and `right = N - 1`.
5. Run the standard two-pointer check for `target = -nums[i]`.

6. When a matching triplet is found, increment `left` and decrement `right`, skipping any identical adjacent values.

```
def three_sum(nums: list[int]) → list[list[int]]:
    nums.sort() # O(N log N)
    n = len(nums)
    triplets = []

    for i in range(n - 2):
        # Optimization: If the smallest number > 0, three numbers cannot sum to 0
        if nums[i] > 0:
            break

        # Avoid duplicate triplets for the first element
        if i > 0 and nums[i] == nums[i - 1]:
            continue

        left, right = i + 1, n - 1
        target = -nums[i]

        while left < right:
            two_sum = nums[left] + nums[right]
            if two_sum == target:
                triplets.append([nums[i], nums[left], nums[right]])
                left += 1
                right -= 1

            # Skip duplicate elements for second and third positions
            while left < right and nums[left] == nums[left - 1]:
                left += 1
            while left < right and nums[right] == nums[right + 1]:
                right -= 1

            elif two_sum < target:
                left += 1
            else:
                right -= 1

    return triplets

# Test cases
print(three_sum([-1, 0, 1, 2, -1, -4])) # [[-1, -1, 2], [-1, 0, 1]]
print(three_sum([0, 1, 1]))           # []
print(three_sum([0, 0, 0]))           # [[0, 0, 0]]
```

- **Time Complexity:** $O(N^2)$ (Sorting takes $O(N \log N)$, outer loop runs N times, inner two-pointer scan takes $O(N)$).
- **Space Complexity:** $O(1)$ auxiliary space (excluding sorting stack space, typically $O(\log N)$).
- **Edge Cases:** Array with all zeroes, all positive, all negative, array with identical duplicates.

7. How It Works Internally: Invariant Preservation

The mathematical foundation of two pointers is **search space reduction**. – In an unsorted array, there are $\binom{N}{2} = \frac{N(N-1)}{2}$ possible pairs. – By sorting the array, we establish a strict monotonic invariant: if $\text{nums}[\text{left}] + \text{nums}[\text{right}] > \text{target}$, then *no other element*

paired with `nums[right]` could ever equal target, because all other available elements are $\geq \text{nums}[\text{left}]$. - Therefore, `right` can be safely discarded forever in $O(1)$ step. Each comparison eliminates an entire row or column from the candidate matrix.

8. Common Mistakes

- Failing to Skip Duplicates in 3-Sum:** Not skipping duplicate elements after recording a triplet results in duplicate entries in the output list.
- Off-by-One While Skipping Duplicates:** Incrementing `left` without checking `left < right` can cause index out-of-bounds exceptions.
- Applying Converging Pointers to Unsorted Arrays:** Converging two pointers *only* works if the search space exhibits monotonicity (e.g., sorted array or geometric boundary).

9. Performance / Complexity Matrix

Problem	Naive Complexity	Two-Pointer Time	Auxiliary Space
Two Sum (Sorted)	$O(N^2)$	$O(N)$	$O(1)$
3-Sum	$O(N^3)$	$O(N^2)$	$O(1)$
Container With Most Water	$O(N^2)$	$O(N)$	$O(1)$
Trapping Rain Water	$O(N^2)$	$O(N)$	$O(1)$
Remove Duplicates from Sorted Array	$O(N^2)$	$O(N)$	$O(1)$

10. Interview Questions (Easy $\rightarrow$ Medium $\rightarrow$ Hard)

Easy

- Q:** Explain how to detect a cycle in a singly linked list using fast and slow pointers (Floyd's Tortoise and Hare algorithm).

Medium

- Q:** Given an array representing vertical lines, find two lines that together with the x-axis form a container containing the most water (*Container With Most Water*).

Hard

- Q:** Given N non-negative integers representing an elevation map where the width of each bar is 1, compute how much water it can trap after raining using $O(1)$ auxiliary space.

11. Follow-up Questions from Interviewer

- "In Container With Most Water, why is it mathematically safe to advance the pointer pointing to the shorter line? Could advancing the taller line ever yield a larger area?"
- "Can you extend 3-Sum to 4-Sum and K -Sum? What is the generalized time complexity for K -Sum?" (Answer: $O(N^{K-1})$ time).

12. Model Answer: Container With Most Water Greediness

30-SECOND INTERVIEW ANSWER (EXECUTIVE PITCH)

Interviewer: "Why do we always move the pointer pointing to the shorter line?"

Model Answer: "The area of water between two lines at indices L and R is given by: $\text{Area} = \min(H[L], H[R]) \times (R - L)$

The width $(R - L)$ starts at its maximum possible value and strictly decreases by 1 on every step.

The height of the water is strictly capped by the **shorter line**.

- If we were to move the **taller line** inwards, the width decreases, but the maximum possible height is still constrained by the unchanged shorter line (or an even shorter one). Therefore, moving the taller line can **never** increase the area.
- The *only* hope of finding a larger area despite the decreasing width is to find a significantly taller line to replace the current limiting bottleneck. Thus, we must advance the pointer pointing to the shorter line.

This greedy choice guarantees that no potential maximum area is ever skipped, enabling an optimal $O(N)$ single-pass solution with $O(1)$ auxiliary space."

```
def max_area(height: list[int]) -> int:
    left, right = 0, len(height) - 1
    max_water = 0

    while left < right:
        width = right - left
        current_water = min(height[left], height[right]) * width
        max_water = max(max_water, current_water)

        if height[left] < height[right]:
            left += 1
        else:
            right -= 1

    return max_water
```

13. Practical Exercise

Implement **Trapping Rain Water** using two pointers: Maintain `left_max` and `right_max`. If `height[left] < height[right]`, process `left`: if `height[left] ≥ left_max`, update `left_max`; otherwise, add `left_max - height[left]` to total water. Then increment `left`. Handle `right`

symmetrically.

14. Quick Revision

- Converging pointers require monotonic sorted arrays.
- Fast and slow pointers find cycle lengths and midpoints of linked lists.
- 3-Sum = Sort ($O(N \log N)$) + Outer loop (N) + Two pointers (N) = $O(N^2)$.
- Always advance the bottleneck pointer in area maximization problems.

15. Interview Checklist

- ☐ Writes 3-Sum cleanly without duplicate triplets.
- ☐ Proves mathematically why the shorter line moves in Container With Most Water.
- ☐ Solves Trapping Rain Water in $O(N)$ time and $O(1)$ auxiliary space.
- ☐ Can explain fast and slow pointer cycle detection.



CHAPTER 04 SUMMARY & REVISION CHECKLIST



WHAT TO REMEMBER

- Zero-indexed array addressing.
- Sorting fundamentals (understanding that pre-sorting an array takes $O(N \log N)$ and often unlocks $O(N)$ two-pointer t
- Time Complexity: $O(N^2)$ (Sorting takes $O(N \log N)$, outer loop runs N times, inner two-pointer scan takes $O(N)$).



CANDIDATE SELF-CHECK

- ☐ I understand the fundamental mechanism
- ☐ I can explain it in natural Hinglish
- ☐ I can write or debug the core code/query
- ☐ I can defend trade-offs and edge cases

CHAPTER 05

Sliding Window Algorithmic Patterns

PRIORITY

● High Priority

DIFFICULTY

Intermediate to Advanced

EST. TIME

35–45 min

FREQUENCY

Very High (85%+)

Sliding Window Algorithmic Patterns

1. Why This Matters

Sliding window algorithms are among the most frequently tested algorithmic techniques at IDFC FIRST Bank and across the FinTech industry. They appear in problems involving contiguous subarrays, string substrings, rolling rate limiters, burst traffic detection, and streaming transaction analytics. The sliding window reduces $O(N^2)$ brute-force window evaluations into optimal $O(N)$ linear time.

2. Prerequisites

- Two-pointer basics.
- Hash map and hash set frequency tracking.

3. Concept

Types of Sliding Windows

- Fixed-Size Window:** The window length K is constant. As the right pointer advances, the left pointer advances simultaneously (`left = right - K + 1`), adding the incoming element and removing the outgoing element in $O(1)$.
- Dynamic / Variable-Size Window:** The window expands by advancing the `right` pointer to satisfy an invariant, and contracts by advancing the `left` pointer when the invariant is violated (e.g., finding the shortest or longest subarray meeting a condition).

```
flowchart LR
    subgraph Window ["Dynamic Window Expansion & Contraction"]
        direction TB
        w1["1. Expand: right += 1 (Add nums[right] to window state)"]
        w2["2. Check: Is window invalid?"]
        w3["3. Shrink: while invalid: left += 1 (Remove nums[left])"]
        w4["4. Record: Update global answer (max_len or min_len)"]
        w1 --> w2 --> w3 --> w4
    end
    end
```

4. Simple Example: Maximum Sum Subarray of Fixed Size K

```
def max_sub_array_fixed(nums: list[int], k: int) → int:
    """
    Finds maximum sum of any contiguous subarray of size k.
    Complexity: O(N) Time, O(1) Space.
    """
    if len(nums) < k:
        return 0

    # Initial window sum
    window_sum = sum(nums[:k])
    max_sum = window_sum

    # Slide window
    for right in range(k, len(nums)):
        window_sum += nums[right] - nums[right - k] # Add incoming, subtract outgoing
        max_sum = max(max_sum, window_sum)

    return max_sum
```

5. Real-World Banking Example: Rolling Rate Limiter & Transaction Spikes

To prevent card testing attacks, an API gateway monitors incoming payment attempts: "Has a merchant terminal initiated more than 100 requests in any rolling 60-second window?" - A sliding window tracks request timestamps. As time moves forward, timestamps older than 60 seconds expire from the left edge of the window, maintaining continuous real-time throughput metrics without resetting on artificial minute boundaries.

6. Code: Longest Substring Without Repeating Characters

Problem Statement

Given a string `s`, find the length of the **longest substring** without repeating characters.

Clarifying Questions

1. What characters can the string contain? English letters, digits, symbols, and spaces.
2. Can the input be empty? Yes, empty string returns 0.

Brute Force Approach

Check every possible substring ($O(N^2)$), and for each check whether all characters are unique using a set ($O(N)$). - **Time Complexity:** $O(N^3)$ - **Space Complexity:** $O(\min(N, M))$ where M is alphabet size.

Optimized Approach: Dynamic Sliding Window with Last Seen Index Map

Instead of shrinking the window one character at a time, store the **most recent index** where each character was seen. When a duplicate character is encountered, jump `left` directly past the previous instance of that character: $\text{left} = \max(\text{left}, \text{last_seen}[\text{char}] + 1)$

```
def length_of_longest_substring(s: str) → int:
    last_seen = {} # char → index
    left = 0
    max_len = 0

    for right, char in enumerate(s):
        # If char was seen and is within the current window
        if char in last_seen and last_seen[char] ≥ left:
            left = last_seen[char] + 1

        last_seen[char] = right
        max_len = max(max_len, right - left + 1)

    return max_len

# Test cases
print(length_of_longest_substring("abcabcbb")) # 3 ("abc")
print(length_of_longest_substring("bbbbb"))    # 1 ("b")
print(length_of_longest_substring("pwwkew"))    # 3 ("wke")
print(length_of_longest_substring(""))          # 0
```

- **Time Complexity:** $O(N)$ (Each character is processed at most once by `right`).
- **Space Complexity:** $O(\min(N, M))$ where M is the character set size (e.g., 128 for ASCII).
- **Edge Cases:** Empty string, string with all unique characters, string with identical characters, strings with spaces and punctuation.

7. How It Works Internally: Amortized Analysis of Sliding Windows

Junior engineers often mistakenly evaluate dynamic sliding windows as $O(N^2)$ because they see a `while` loop nested inside a `for` loop:

```
for right in range(n):
    while window_is_invalid():
        left += 1
```

Why it is strictly $O(N)$ Amortized: – The `right` pointer increments exactly N times. – The `left` pointer only ever moves forward; it *never* decrements or resets. – Therefore, across the entire execution of the algorithm, `left` increments at most N times. – Total pointer operations $= N + N = 2N$ implies $O(N)$ time.

8. Common Mistakes

1. **Forgetting `last_seen[char] ≥ left`:** In the index jump optimization, if a character was seen at index 1, but `left` has already advanced to index 5, forgetting this check pulls `left` *backwards*, corrupting the window!
2. **Incorrect Window Size Calculation:** Window length is always $\text{right} - \text{left} + 1$ (inclusive), not $\text{right} - \text{left}$.

3. **Attempting Sliding Window on Negative Number Arrays:** Finding a subarray summing to K cannot use a sliding window when negative numbers are present, because expanding the window can *decrease* the sum.

9. Performance / Complexity Matrix

Problem	Naive Time	Sliding Window Time	Auxiliary Space
Max Sum Subarray of Size K	$O(N \cdot K)$	$O(N)$	$O(1)$
Longest Substring Without Repeating Characters	$O(N^3)$	$O(N)$	$O(M)$
Minimum Size Subarray Sum	$O(N^2)$	$O(N)$	$O(1)$
Sliding Window Maximum	$O(N \cdot K)$	$O(N)$ (Monotonic Deque)	$O(K)$

10. Interview Questions (Easy to Medium to Hard)

Easy

- Q:** Explain how a fixed-size sliding window updates its running sum in $O(1)$ time.

Medium

- Q:** Given an array of positive integers `nums` and a positive integer `target`, find the minimal length of a contiguous subarray of which the sum is greater than or equal to `target`. If there is no such subarray, return 0.

Hard

- Q:** You are given an array of integers `nums`, there is a sliding window of size `k` which is moving from the very left of the array to the very right. Return the max sliding window. Solve it in $O(N)$ time using a Monotonic Deque.

11. Follow-up Questions from Interviewer

- "In the Monotonic Deque solution for Sliding Window Maximum, why do we store indices in the deque rather than values?" (Answer: Storing indices allows us to verify whether the oldest element has fallen out of the sliding window: `deque[0] ≤ right - k`).
- "How does a leaky bucket rate limiter algorithm differ conceptually from a sliding window log rate limiter?"

12. Model Answer: Minimum Size Subarray Sum

30-SECOND INTERVIEW ANSWER (EXECUTIVE PITCH)

Interviewer: "Walk me through how you solve Minimum Size Subarray Sum in $O(N)$."

Model Answer: "Because all numbers in `nums` are strictly positive, the cumulative sum within a window increases monotonically as we expand `right`, and decreases monotonically as we shrink `left`. This satisfies the sliding window property."

1. We maintain a running `current_sum` and initialize `min_len = infinity`.
2. We expand the window by adding `nums[right]` to `current_sum`.
3. Once `current_sum ≥ target`, we attempt to minimize the window: while `current_sum ≥ target`, we update `min_len = min(min_len, right - left + 1)`, subtract `nums[left]` from `current_sum`, and increment `left`.
4. Because each element enters the window once via `right` and leaves at most once via `left`, the total operations are bounded by $2N$. Time complexity is strictly $O(N)$, and auxiliary space is $O(1)$.

```
def min_sub_array_len(target: int, nums: list[int]) → int:
    left = 0
    current_sum = 0
    min_len = float("inf")

    for right in range(len(nums)):
        current_sum += nums[right]

        while current_sum ≥ target:
            min_len = min(min_len, right - left + 1)
            current_sum -= nums[left]
            left += 1

    return min_len if min_len ≠ float("inf") else 0
```

13. Practical Exercise

Implement **Sliding Window Maximum** in $O(N)$ using Python's `collections.deque`: Maintain a deque of indices where values are strictly in descending order. Before appending `right`, pop all smaller elements from the back of the deque. Pop the front element if its index is $\text{right} - k$.

14. Quick Revision

- Fixed window: `window_sum += incoming - outgoing`.
- Dynamic window: expand with `right`, shrink with `left` while condition holds.
- Complexity is $O(N)$ amortized because each element is visited at most twice.
- Window size $= \text{right} - \text{left} + 1$.
- Negative numbers invalidate sliding window monotonicity; use Prefix Sum + Hash Map instead.

15. Interview Checklist

- [] Writes Longest Substring Without Repeating Characters with index jumps.
- [] Proves $O(N)$ amortized complexity for nested while loop.
- [] Remembers that sliding window requires positive numbers for sum thresholds.
- [] Understands Monotonic Deque for Sliding Window Maximum.



CHAPTER 05 SUMMARY & REVISION CHECKLIST



WHAT TO REMEMBER

- Two-pointer basics.
- Hash map and hash set frequency tracking.
- A sliding window tracks request timestamps. As time moves forward, timestamps older than 60 seconds expire from the left



CANDIDATE SELF-CHECK

- ☐ I understand the fundamental mechanism
- ☐ I can explain it in natural Hinglish
- ☐ I can write or debug the core code/query
- ☐ I can defend trade-offs and edge cases

CHAPTER 06

Python for Coding Interviews & DSA Mastery

PRIORITY ● High Priority	DIFFICULTY Intermediate to Advanced	EST. TIME 35–45 min	FREQUENCY Very High (85%+)
--	---	---	--

Python for Coding Interviews & DSA Mastery

1. Why This Matters

Coding interviews mein speed aur clean code sabse zyada count hota hai. Python ki standard library mein aisi built-in data structures hain (`deque` , `heapq` , `Counter` , `defaultdict` , `bisect`) jo aapko 50 lines ka boilerplate likhne ke bajaye 2 lines mein optimal logic implement karne deti hain. Lekin agar aapko inka internal time complexity aur edge cases nahi pata (jaise `list.pop(0)` ka $O(N)$ hona), toh interviewer aapko complexity round mein disqualify kar dega.

2. Python DSA Decision Tree: "Kab Kya Use Karein?"

```
flowchart TD
    Problem["Algorithmic Requirement"] --> R1["Unique elements or lookup in O(1)?"]
    R1 -->|Yes| S["Use set or dict"]
    R1 -->|No| R2["FIFO Queue or BFS?"]
    R2 -->|Yes| DQ["Use collections.deque (popleft O(1))"]
    R2 -->|No| R3["Top K elements / Min / Max in O(log N)?"]
    R3 -->|Yes| HQ["Use heapq (Min-Heap / Negated Max-Heap)"]
    R3 -->|No| R4["Frequency counts or grouped lists?"]
    R4 -->|Yes| C["Use collections.Counter or defaultdict"]
    R4 -->|No| R5["Binary Search on sorted data?"]
    R5 -->|Yes| B["Use bisect.bisect_left / bisect_right"]
    R5 -->|No| L["Use standard list"]
```

3. Data Structure Deep-Dive & Complexity Matrix

1. `list` — Dynamic Array

- **When to use:** Index-based access, sequential scans, stacks (LIFO via `append()` and `pop()`).
- **Under the hood:** Contiguous array of pointers (`PyObject**`). Over-allocates capacity geometrically.

Method	Time Complexity	Notes & Gotchas
<code>arr[i]</code>	$O(1)$	Direct pointer arithmetic in memory
<code>arr.append(x)</code>	$O(1)$ amortized	Geometric resizing
<code>arr.pop()</code>	$O(1)$	Pops from end (Stack LIFO)
<code>arr.pop(0)</code>	$O(N)$ (DANGEROUS)	Shifts all $N-1$ pointers left in RAM!
<code>arr.insert(0, x)</code>	$O(N)$	Shifts all elements right
<code>x in arr</code>	$O(N)$	Linear scan

2. `collections.deque` — Double-Ended Queue

- **When to use:** BFS (Breadth-First Search), Sliding Window Maximum, FIFO queues.
- **Under the hood:** Doubly-linked list of 64-element block buffers.

```
from collections import deque

queue = deque([10, 20, 30])
queue.append(40)      # O(1) Push Right
queue.appendleft(0)   # O(1) Push Left
first = queue.popleft() # O(1) Pop Left (NEVER use list.pop(0)!)
last = queue.pop()    # O(1) Pop Right
```

3. `heapq` — Priority Queue (Min-Heap)

- **When to use:** "Top K frequent elements", "Kth largest element", "Median of a data stream", Dijkstra's algorithm.
- **Under the hood:** Binary Min-Heap implemented over a standard Python list. Smallest element is always at `heap[0]`.

```
import heapq

# Min-Heap (Default)
min_heap = []
heapq.heappush(min_heap, 25) # O(log N)
heapq.heappush(min_heap, 5)
heapq.heappush(min_heap, 50)
smallest = heapq.heappop(min_heap) # 5 in O(log N)

# Max-Heap Pattern: Negate values upon push and pop!
max_heap = []
heapq.heappush(max_heap, -25)
heapq.heappush(max_heap, -5)
heapq.heappush(max_heap, -50)
largest = -heapq.heappop(max_heap) # 50

# In-Place Heapify: Converts an existing list into a heap in O(N) linear time!
nums = [15, 3, 20, 1, 9]
heapq.heapify(nums) # O(N) time (Faster than pushing N times O(N log N)!)
```

4. `collections.Counter` & `collections.defaultdict`

- **When to use:** Character counts, anagrams, grouping items by key without `KeyError`.

```
from collections import Counter, defaultdict

# 1. Counter: Instant frequency map
counts = Counter("idfcfirstbank")
print(counts["i"]) # 2
print(counts.most_common(2)) # [('i', 2), ('f', 1)]

# 2. defaultdict: Eliminates 'if key not in dict' boilerplate
# Graph adjacency list representation:
adj_list = defaultdict(list)
edges = [(1, 2), (1, 3), (2, 4)]
for u, v in edges:
    adj_list[u].append(v)
# adj_list[1] → [2, 3]; adj_list[5] → [] (No KeyError!)
```

5. `bisect` — Binary Search Module

- **When to use:** Finding insertion point in sorted arrays, range query boundaries in $O(\log N)$.

```
import bisect

sorted_rates = [3.5, 5.0, 7.25, 8.5, 9.1]

# bisect_left: First index where val can be inserted maintaining sort order (≥ val)
idx_left = bisect.bisect_left(sorted_rates, 7.25) # Index 2

# bisect_right: Index after any existing equal elements (> val)
idx_right = bisect.bisect_right(sorted_rates, 7.25) # Index 3
```

4. Custom Sorting & Comparator Keys

In technical interviews, sorting complex objects or breaking ties is very common:

```
# Problem: Sort transactions by:
# 1. Status: 'FAILED' first, then 'PENDING', then 'SUCCESS'
# 2. Amount: Descending (Higher amounts first)
# 3. Reference ID: Ascending (Lexicographical)

transactions = [
    {"ref": "TX1", "amount": 5000, "status": "SUCCESS"},
    {"ref": "TX2", "amount": 12000, "status": "FAILED"},
    {"ref": "TX3", "amount": 8000, "status": "FAILED"},
    {"ref": "TX4", "amount": 5000, "status": "SUCCESS"},
]

status_priority = {"FAILED": 0, "PENDING": 1, "SUCCESS": 2}

# Custom key function returning a tuple:
sorted_txns = sorted(
    transactions,
    key=lambda x: (status_priority[x["status"]], -x["amount"], x["ref"])
)
# Result: TX2 (12k FAILED), then TX3 (8k FAILED), then TX1 (5k SUCCESS), then TX4
```

5. Recursion, BFS & DFS Templates in Python

1. BFS Template (Iterative with `deque`)

```
def bfs_shortest_path(graph, start, target):
    queue = deque([(start, 0)]) # (node, distance)
    visited = {start}

    while queue:
        node, dist = queue.popleft() # 0(1)
        if node == target:
            return dist

        for neighbor in graph[node]:
            if neighbor not in visited:
                visited.add(neighbor)
                queue.append((neighbor, dist + 1))

    return -1 # Target unreachable
```

2. DFS Template (Recursive with Visited Set)

```
def dfs_traversal(graph, node, visited=None):
    if visited is None:
        visited = set()

    visited.add(node)
    # Process current node
    for neighbor in graph[node]:
        if neighbor not in visited:
            dfs_traversal(graph, neighbor, visited)

    return visited
```

6. Top 5 Python DSA Interview Traps

- Using `list.pop(0)` in BFS:** Turns $O(V + E)$ into quadratic $O(V^2)$ because each pop shifts entire list memory. Always use `collections.deque.popleft()`.
- Mutating list while iterating:** `for x in arr: if condition: arr.remove(x)` skips elements because indices shift dynamically. Create a new list via comprehension or iterate backwards.
- Using 2D lists with multiplication:** `matrix = [[0] * 3] * 3` creates 3 references to the **same single row**! Modifying `matrix[0][0] = 1` mutates all 3 rows. Always use: `matrix = [[0] * cols for _ in range(rows)]`.
- Neglecting recursion depth limit:** Python default recursion depth is 1,000. In deep graph/tree problems, `sys.setrecursionlimit(200000)` is required or convert recursion to iterative stack.
- Float comparison precision:** `0.1 + 0.2 == 0.3` evaluates to `False`. Use `math.isclose(a, b)` for floating-point equality.



CHAPTER 06 SUMMARY & REVISION CHECKLIST



WHAT TO REMEMBER

- When to use: Index-based access, sequential scans, stacks (LIFO via `append()` and `pop()`).
- Under the hood: Contiguous array of pointers (`PyObject``). Over-allocates capacity geometrically.
- When to use: BFS (Breadth-First Search), Sliding Window Maximum, FIFO queues.



CANDIDATE SELF-CHECK

- ☐ I understand the fundamental mechanism
- ☐ I can explain it in natural Hinglish
- ☐ I can write or debug the core code/query
- ☐ I can defend trade-offs and edge cases

CHAPTER 07

DSA Last–Minute Algorithmic Patterns Cheatsheet

PRIORITY ● High Priority	DIFFICULTY Intermediate to Advanced	EST. TIME 35–45 min	FREQUENCY Very High (85%+)
--	---	---	--

DSA Last–Minute Algorithmic Patterns Cheatsheet

The Top 14 Algorithmic Patterns Quick Reference

Pattern	Trigger in Problem Statement	Core Data Structure	Typical Complexity
1. Frequency Hash Map	"Find duplicates", "Frequency of elements", "Anagrams"	<code>collections.Counter</code> / <code>dict</code>	$O(N)$ Time, $O(N)$ Space
2. Two Pointers	"Sorted array", "Find pair summing to target", "Palindromes"	Two integer indices (<code>L</code> , <code>R</code>)	$O(N)$ Time, $O(1)$ Space
3. Sliding Window	"Contiguous subarray/substring", "Longest/Shortest with condition"	Two indices + Hash Map	$O(N)$ Time, $O(K)$ Space
4. Fast & Slow Pointers	"Cycle detection in linked list", "Find middle of linked list"	<code>slow</code> , <code>fast</code> pointers	$O(N)$ Time, $O(1)$ Space
5. Prefix Sum	"Subarray sum equals K", "Range sum queries"	Array / Running sum	$O(N)$ Time, $O(1)$ query
6. Monotonic Stack	"Next greater element", "Daily temperatures", "Stock span"	<code>list</code> as stack	$O(N)$ Time, $O(N)$ Space
7. Top K Elements	"Find K largest / K most frequent", "Running median"	<code>heapq</code> (Min-Heap / Max-Heap)	$O(N \log K)$ Time
8. Binary Search	"Sorted input", "Find peak", "Search rotated sorted array", "Min of Max"	<code>left</code> , <code>right</code> , <code>mid</code>	$O(\log N)$ Time, $O(1)$ Space
9. BFS (Level Order)	"Shortest path in unweighted graph/grid", "Level-by-level tree"	<code>collections.deque</code>	$O(V + E)$ Time, $O(V)$ Space
10. DFS (Backtracking)	"Generate all combinations/permutations/subsets", "Islands in grid"	Recursion / Call Stack	$O(2^N)$ or $O(N!)$ Time
11. Merge Intervals	"Overlapping intervals", "Meeting rooms schedule"	Sort by <code>start_time</code>	$O(N \log N)$ Time
12. Dynamic Programming	"Overlapping subproblems", "Max/Min path", "Number of ways"	1D / 2D table or memo dict	$O(N)$ or $O(N \times M)$
13. Greedy	"Interval scheduling", "Jump game", "Local optimum yields global"	Priority Queue or Sort	$O(N \log N)$ Time
14. Union-Find (DSU)	"Connected components", "Detect cycle in undirected graph"	<code>parent</code> and <code>rank</code> arrays	$O(\alpha(N))$ approx $O(1)$

Python High-Yield Syntax Snippets

```
# 1. Monotonic Deque for O(1) pops and appends:
from collections import deque
q = deque()
q.append(1)
q.popleft() # O(1) vs list.pop(0) which is O(N)!

# 2. Min-Heap & Max-Heap:
import heapq
min_h = []
heapq.heappush(min_h, 10)
val = heapq.heappop(min_h)
# Max-heap: push negative values
heapq.heappush(min_h, -10)
val = -heapq.heappop(min_h)

# 3. Safe Binary Search Mid Calculation:
mid = left + (right - left) // 2

# 4. In-Place Reversal:
def reverse_range(arr, l, r):
    while l < r:
        arr[l], arr[r] = arr[r], arr[l]
        l += 1
        r -= 1

# 5. Defaultdict without KeyError:
from collections import defaultdict
graph = defaultdict(list)
counts = defaultdict(int)
```

CHAPTER 07 SUMMARY & REVISION CHECKLIST

WHAT TO REMEMBER

- Core fundamentals and architectural mechanics of DSA Last-Minute Algorithmic Patterns Cheatsheet
- Production banking scenarios and transaction reliability
- High-frequency technical interview questions and follow-ups

CANDIDATE SELF-CHECK

- ☐ I understand the fundamental mechanism
- ☐ I can explain it in natural Hinglish
- ☐ I can write or debug the core code/query
- ☐ I can defend trade-offs and edge cases